Minimizing the Sum of Piecewise Linear Convex Functions

Yu Cong

October 4, 2024

Plan

$\min \sum f_i(a_i \cdot x - b_i)$

Problem

Given n piecewise linear convex functions $f_1, ..., f_n : \mathbb{R} \to \mathbb{R}$ of total m breakpoints, and n linear functions $a_i \cdot x - b_i : \mathbb{R}^d \to \mathbb{R}$, find $\min_x \sum_i f_i(a_i \cdot x - b_i)$.

(a) A 1D piecewise linear function with 4 line segments and 3 breakpoints

(b) A 2D piecewise concave function

$f_i(a_i \cdot x - b_i) : \mathbb{R}^d \to \mathbb{R}$ is also piecewise linear convex.

General piecewise linear convex function in $\mathbb{R}^d$

Definition (piecewise linear convex function in $\mathbb{R}^d$)

$$g(x) = \max\{a_1^T x + b_1, \ldots, a_L^T x + b_L\}$$

Every piecewise linear convex function in $\mathbb{R}^d$ can be expressed in this form.[1]

However, observe that in our problem the piecewise linear convex function is not that general. It is a composition of a linear mapping and an 1D piecewise linear convex function.

---

[1] S.P. Boyd, L. Vandenberghe, **Convex optimization**, Cambridge University Press, Cambridge, UK ; New York, 2004.

$f \circ l \not\equiv g$

Proof.

Consider a piecewise linear convex function $g : \mathbb{R}^2 \to \mathbb{R}$. g can be viewed as the maximum of a set of planes in $\mathbb{R}^3$.

Consider a series of points $P = \{p_1, p_2, ..., p_k\}$ on the 2D plane. After applying the linear mapping to P, we will get a sequence of numbers(points in 1D) $P' = \{p'_1, p'_2, ..., p'_k\}$. We assume that P' is non-decreasing. Note that the value of g on P' is always unimodal since g is convex. However, the value of f on P may not be unimodal. Thus the composition of a linear mapping and a pwl convex function in 1D is not equivalent to pwl convex functions in high dimensions. $\square$

A linear time algorithm I

Problem

Given n piecewise linear convex functions $f_1, ..., f_n : \mathbb{R} \to \mathbb{R}$ of total m breakpoints, and n linear functions $a_i \cdot x - b_i : \mathbb{R}^d \to \mathbb{R}$, find $\min_x \sum_i f_i(a_i \cdot x - b_i)$.

This can be solve in $O(2^{2^d}(m+n))$ through Megiddo's Low dimension LP algorithm.[2]

Let n_i be the number of line segments in f_i. Note that $\sum_i n_i = m + n$.

We can formulate the optimization problem as the following linear program,

A linear time algorithm II

$$\min \sum_{i=1}^{n} f_i$$
$$s.t. \quad f_i \geq \alpha_j(a_i \cdot x - b_i) - \beta_j \quad \forall i \in [n], \forall j$$

where $\alpha_j x - \beta_j$ is the j'th line segment on f_i.
There will be $m + n$ constraints in total.

---

[2]Nimrod Megiddo. Linear programming in linear time when the dimension is fixed. J. ACM, 31(1):114–127, jan 1984.

Megiddo's algorithm I

The dimension d (in our problem, the dimension of x) is small while the number of constraints are huge. We need only d linearly independent tight constraints to identify the optimal solution x^*. Thus most of the constraints are useless.

For one constraint, how can we know where does x^* locate with respect to it?
Through inquiries. Let $a \cdot x \leq b$ be the constraint. Define 3 hyperplanes, $a \cdot x = c$ where $c \in \{b, b - \varepsilon, b + \varepsilon\}$. Now solve three $d - 1$ dimension linear programming. The largest of the three objective functions tells us where x^* lies with respect to the hyperplane.

Megiddo's algorithm II

Finding the optimal solution x^* is therefore equivalent to the following problem,

Problem (Multidimensional Search Problem)

Suppose that there exists a point x^* which is not known to us, but there is a oracle that can tell the position of x^* relative to any hyperplane in $\mathbb{R}^d$. Given n hyperplanes, we want to know the position of x^* relative to each of them.

What about 1 dimension search? A fastest way will be using the linear time median algorithm. We can find the median of n numbers and call the oracle to compare the median with x^*. Thus with $O(n)$ time median finding and one oracle call, we find the relative position of $n/2$ elements relative to x^*.

Megiddo's algorithm III

If we can do similar things in $\mathbb{R}^d$, i.e., there is a method which makes $A(d)$ oracle calls and determines at least $B(d)$ fraction of relative positions, then we can apply this method $\log_{\frac{1}{1-B(d)}} n$ times to find all relative positions.

Note that in 1 dimension, $A(1) = 1$ and $B(1) = 1/2$ (call oracle to compare x^* and the median). In $\mathbb{R}^d$, our oracle is the recursive inquiry.

A trivial method will be iterating on all hyperplanes and calling the oracle on each one, since there is no *median* of a set of hyperplanes in $\mathbb{R}^d$. The complexity recurrence is

$$T(n, d) = n(3T(n - 1, d - 1) + O(nd))$$

Note that in this setting $A(d) = 1$ and $B(d) = 1/n$.

Megiddo's algorithm IV

Megiddo designed a clever method where $A(d) = 2^{d-1}$ and $B(d) = 2^{-(2^d-1)}$.

Lemma

Given two lines through the origin with slopes of opposite sign, knowing which quadrant x^* lies in allows us to locate it with respect to at least one of the lines.

Megiddo's algorithm V

Let l_H be the intersection of hyperplane H and $x_1 x_2$ plane.
Compute a partition $S_1 \sqcup S_2 = \mathcal{H}$. $H \in S_1$ iff l_H has positive slope.
Otherwise $l_H \in S_2$. We further assume that $|S_1| = |S_2| = n/2$.

Now we have $n/2$ pairs (H_1, H_2), where $H_i \in S_i$. Let l_i be the intersection of H_i and $x_1 x_2$ plane. Let H_{x_i} be the linear combination of H_1 and H_2 s.t. x_i is eliminated.

By the previous lemma, calling oracle on l_{x_1} and l_{x_2} locate x^* with respect to at least one of H_1 and H_2.

Megiddo's algorithm VI

Input: S_1, S_2 and the pairs.

1. recursively locate x^* respect to $B(d-1)n/2$ hyperplanes(H_{x_i}) with $A(d-1)$ oracle calls in S_1.

2. locate with respect to a $B(d-1)$-fraction of corresponding paired hyperplanes in S_2.

3. There are still $(1 - B(d-1)^2)/2$-fraction of hyperplanes for which we do not know the relative position with x^*. Run this algorithm on these hyperplanes.

This gives the recurrence

$$T(n, d) \leq 3 \cdot 2^{d-1} T(n, d-1) + T((1 - 2^{1-2^d})n, d) + O(nd)$$

with solution $T(n, d) = O(2^{2^d} n)$.

Zemel's conversion

Note that our linear program is not exactly the same as low
dimension LP.

Other algorithms for fixed dimension LP

simplex method	det.	$O(n/d)^{d/2+O(1)}$
Megiddo [24]	det.	$2^{O(2^d)}n$
Clarkson [9]/Dyer [14]	det.	$3^{d^2}n$
Dyer and Frieze [15]	rand.	$O(d)^{3d}(\log d)^d n$
Clarkson [10]	rand.	$d^2n + O(d)^{d/2+O(1)}\log n + d^4\sqrt{n}\log n$
Seidel [26]	rand.	$d!n$
Kalai [19]/Matoušek, Sharir, and Welzl [23]	rand.	$\min\{d^2 2^d n,\ e^{2\sqrt{d\ln(n/\sqrt{d})}+O(\sqrt{d}+\log n)}\}$
combination of [10] and [19, 23]	rand.	$d^2n + 2^{O(\sqrt{d\log d})}$
Hansen and Zwick [18]	rand.	$2^{O(\sqrt{d\log((n-d)/d)})}n$
Agarwal, Sharir, and Toledo [4]	det.	$O(d)^{10d}(\log d)^{2d}n$
Chazelle and Matoušek [8]	det.	$O(d)^{7d}(\log d)^d n$
Brönnimann, Chazelle, and Matoušek [5]	det.	$O(d)^{5d}(\log d)^d n$
~~this paper~~ Chan	det.	$O(d)^{d/2}(\log d)^{3d}n$

Figure: Algorithms for LP in low dimensions [3]

Question. Can we use faster fixed dimension LP algorithms to get better complexity?

[3]table stolen from https://dl.acm.org/doi/10.1145/3155312

LP type problem

Aggregate the pwl convex functions