

Minimizing the Sum of Piecewise Linear Convex Functions

Yu Cong

October 6, 2024


Plan

1. Problems & Definitions
2. Properties
3. LP in Low Dimensions
4. Possible Improvements


$$\min \sum f_i(a_i \cdot x - b_i)$$

Problem

Given n piecewise linear convex functions $f_1, \dots, f_n : \mathbb{R} \rightarrow \mathbb{R}$ of total m breakpoints, and n linear functions $a_i \cdot x - b_i : \mathbb{R}^d \rightarrow \mathbb{R}$, find $\min_x \sum_i f_i(a_i \cdot x - b_i)$.



(a) A 1D pwl function with 4 line segments and 3 breakpoints



(b) A 2D pwl concave function

$f_i(a_i \cdot x - b_i) : \mathbb{R}^d \rightarrow \mathbb{R}$ is also piecewise linear convex.

General piecewise linear convex function in \mathbb{R}^d

Definition (piecewise linear convex function in \mathbb{R}^d)

$$g(x) = \max\{a_1^T x + b_1, \dots, a_L^T x + b_L\}$$

Every piecewise linear convex function in \mathbb{R}^d can be expressed in this form.¹

However, observe that in our problem the piecewise linear convex function is not that general. It is a composition of a linear mapping and an 1D piecewise linear convex function.

¹S.P. Boyd, L. Vandenberghe, **Convex optimization**, Cambridge University Press, Cambridge, UK ; New York, 2004.

$$f \circ l \not\equiv g$$

Proof.

Consider a piecewise linear convex function $g : \mathbb{R}^2 \rightarrow \mathbb{R}$. g can be viewed as the maximum of a set of planes in \mathbb{R}^3 .

Consider a series of points $P = \{p_1, p_2, \dots, p_k\}$ on the 2D plane. After applying the linear mapping to P , we will get a sequence of numbers (points in 1D) $P' = \{p'_1, p'_2, \dots, p'_k\}$. We assume that P' is non-decreasing. Note that the value of g on P' is always unimodal since g is convex. However, the value of f on P may not be unimodal. Thus the composition of a linear mapping and a pwl convex function in 1D is not equivalent to pwl convex functions in high dimensions. \square

A linear time algorithm I

Problem

Given n piecewise linear convex functions $f_1, \dots, f_n : \mathbb{R} \rightarrow \mathbb{R}$ of total m breakpoints, and n linear functions $a_i \cdot x - b_i : \mathbb{R}^d \rightarrow \mathbb{R}$, find $\min_x \sum_i f_i(a_i \cdot x - b_i)$.

This can be solve in $O(2^{2^d}(m + n))$ through Megiddo's Low dimension LP algorithm.²

Let n_i be the number of line segments in f_i . Note that

$$\sum_i n_i = m + n.$$

We can formulate the optimization problem as the following linear program,

A linear time algorithm II

$$\begin{aligned} \min \quad & \sum_{i=1}^n f_i \\ \text{s.t.} \quad & f_i \geq \alpha_j(a_i \cdot x - b_i) - \beta_j \quad \forall i \in [n], \forall j \end{aligned}$$

where $\alpha_j x - \beta_j$ is the j 'th line segment on f_i .

There will be $m + n$ constraints in total.

²Nimrod Megiddo. Linear programming in linear time when the dimension is fixed. J. ACM, 31(1):114–127, jan 1984.

Megiddo's algorithm I

The dimension d (in our problem, the dimension of x) is small while the number of constraints are huge. We need only d linearly independent tight constraints to identify the optimal solution x^* . Thus most of the constraints are useless.

For one constraint, how can we know where does x^* locate with respect to it?

Through inquiries. Let $a \cdot x \leq b$ be the constraint. Define 3 hyperplanes, $a \cdot x = c$ where $c \in \{b, b - \varepsilon, b + \varepsilon\}$. Now solve three $d - 1$ dimension linear programming. The largest of the three objective functions tells us where x^* lies with respect to the hyperplane.

Megiddo's algorithm II

Finding the optimal solution x^* is therefore equivalent to the following problem,

Problem (Multidimensional Search Problem)

Suppose that there exists a point x^ which is not known to us, but there is a oracle that can tell the position of x^* relative to any hyperplane in \mathbb{R}^d . Given n hyperplanes, we want to know the position of x^* relative to each of them.*

What about 1 dimension search? A fastest way will be using the linear time median algorithm. We can find the median of n numbers and call the oracle to compare the median with x^* . Thus with $O(n)$ time median finding and one oracle call, we find the relative position of $n/2$ elements relative to x^* .

Megiddo's algorithm III

If we can do similar things in \mathbb{R}^d , i.e., there is a method which makes $A(d)$ oracle calls and determines at least $B(d)$ fraction of relative positions, then we can apply this method $\log_{\frac{1}{1-B(d)}} n$ times to find all relative positions.

Note that in 1 dimension, $A(1) = 1$ and $B(1) = 1/2$ (call oracle to compare x^* and the median). In \mathbb{R}^d , our oracle is the recursive inquiry.

A trivial method will be iterating on all hyperplanes and calling the oracle on each one, since there is no *median* of a set of hyperplanes in \mathbb{R}^d . The complexity recurrence is


$$T(n, d) = n(3T(n-1, d-1) + O(nd))$$

Note that in this setting $A(d) = 1$ and $B(d) = 1/n$.

Megiddo's algorithm IV

Megiddo designed a clever method where $A(d) = 2^{d-1}$ and $B(d) = 2^{-(2^d-1)}$.

Lemma




Given two lines through the origin with slopes of opposite sign, knowing which quadrant x^ lies in allows us to locate it with respect to at least one of the lines.*

Megiddo's algorithm V

Let l_H be the intersection of hyperplane H and x_1x_2 plane.

Compute a partition $S_1 \sqcup S_2 = \mathcal{H}$. $H \in S_1$ iff l_H has positive slope.

Otherwise $l_H \in S_2$. We further assume that $|S_1| = |S_2| = n/2$.



Now we have $n/2$ pairs (H_1, H_2) , where $H_i \in S_i$. Let l_i be the intersection of H_i and x_1x_2 plane. Let H_{x_i} be the linear combination of H_1 and H_2 s.t. x_i is eliminated.

By the previous lemma, calling oracle on l_{x_1} and l_{x_2} locate x^* with respect to at least one of H_1 and H_2 .

Megiddo's algorithm VI

Input: S_1, S_2 and the pairs.

- 1 recursively locate x^* respect to $B(d-1)n/2$ hyperplanes(H_{x_i}) with $A(d-1)$ oracle calls in S_1 .
- 2 locate with respect to a $B(d-1)$ -fraction of corresponding paired hyperplanes in S_2 .
- 3 There are still $(1 - B(d-1)^2)/2$ -fraction of hyperplanes for which we do not know the relative position with x^* . Run this algorithm on these hyperplanes.

This gives the recurrence

$$T(n, d) \leq 3 \cdot 2^{d-1} T(n, d-1) + T((1 - 2^{1-2^d})n, d) + O(nd)$$

with solution $T(n, d) = O(2^{2^d} n)$.

Zemel's conversion

Our linear program has *dimension* $n + d$. Zemel showed that this kind of problem can be converted to a linear program of dimension d .

$$\begin{aligned} \min \quad & \sum_{i=1}^n f_i \\ \text{s.t.} \quad & f_i \geq \alpha_j(a_i \cdot x - b_i) - \beta_j \quad \forall i \in [n], \forall j \end{aligned}$$

Here is an intuitive way to understand the conversion. One can think the LP above as a d -dimensional search problem with $n + d$ hyperplanes. However, the oracle is quite different. The oracle takes the unknown x^* and a hyperplane H as input, returns the relative position by computing the minimal f_i .

Other algorithms for fixed dimension LP

simplex method	det.	$O(n/d)^{d/2+O(1)}$
Megiddo [24]	det.	$2^{O(2^d)}n$
Clarkson [9]/Dyer [14]	det.	$3^{d^2}n$
Dyer and Frieze [15]	rand.	$O(d)^{3d}(\log d)^d n$
Clarkson [10]	rand.	$d^2n + O(d)^{d/2+O(1)} \log n + d^4 \sqrt{n} \log n$
Seidel [26]	rand.	$d!n$
Kalai [19]/Matoušek, Sharir, and Welzl [23]	rand.	$\min\{d^2 2^d n, e^{2\sqrt{d \ln(n/\sqrt{d})} + O(\sqrt{d} + \log n)}\}$
combination of [10] and [19, 23]	rand.	$d^2n + 2^{O(\sqrt{d \log d})}$
Hansen and Zwick [18]	rand.	$2^{O(\sqrt{d \log((n-d)/d)})} n$
Agarwal, Sharir, and Toledo [4]	det.	$O(d)^{10d}(\log d)^{2d}n$
Chazelle and Matoušek [8]	det.	$O(d)^{7d}(\log d)^d n$
Brönnimann, Chazelle, and Matoušek [5]	det.	$O(d)^{5d}(\log d)^d n$
this paper Chan	det.	$O(d)^{d/2}(\log d)^{3d}n$

Figure: Algorithms for LP in low dimensions ³

Can we use faster fixed dimension LP algorithms to get better complexity?

³table stolen from <https://dl.acm.org/doi/10.1145/3155312>

LP-type problem I

Algorithms for low dim LP are actually solving a more abstract problem.

Definition (LP-type problem)

Given a set S and a function $f : S \rightarrow \mathbb{R}$. f satisfies two properties:

- Monotonicity: $\forall A \subseteq B \subseteq S, f(A) \leq f(B) \leq f(S)$.
- Locality: $\forall A \subseteq B \subseteq S$ and $\forall x \in S$, if $f(A) = f(B) = f(A \cup \{x\})$, then $f(A) = f(B \cup \{x\})$.

Linear programs(minimization) are LP-type problems.

$B \subseteq S$ is a basis if $\forall B' \subsetneq B, f(B') < f(B)$. A set of 'useful' constraints in a linear program is a basis.

The combinatorial dimension is the size of the largest basis.

If a LP problem has low dimension, then its combinatorial dimension is low. **What about the converse?**

LP-type problem II

$$\begin{aligned} \min \quad & \sum_{i=1}^n f_i \\ \text{s.t.} \quad & f_i \geq \alpha_j(a_i \cdot x - b_i) - \beta_j \quad \forall i \in [n], \forall j \\ & \dots \end{aligned}$$

Does our LP has low combinatorial dimension?

No. A basis contains at least n constraints since otherwise some f_i is unbounded.

Problem

Is it possible to formulate the pwl convex minimization problem as an LP-type problem with low combinatorial dimension?

Aggregate the pwl convex functions

The sum of pwl convex functions are still pwl convex.